

# Brussels Advanced training

Giles Innocent

2024-01-31

# Introduction

- ▶ Building on the basic course.
- ▶ **NOT** lecture:practical.
- ▶ More intro: practical or discussion.
- ▶ You'll get more out of it if you try the exercises **before** looking at the sample code.
- ▶ Nothing Bayesian, or specific to diagnostic tests or latent class analyses (although it is particularly useful here)
- ▶ There is more than one way to code

# Simulation

- ▶ Why simulate?
- ▶ How to simulate
  - ▶ within JAGS
  - ▶ R or equivalent
  - ▶ from an identical model to the analysis model
  - ▶ from a different model to the analysis

## Simulating in R

- ▶ Functions like rbinom, rpois, rnorm, etc.
- ▶ All take a first parameter, n the number of data points you wish to simulate
- ▶ E.G. Hui-Walter paradigm:

```
set.seed(1)
n.sim <- 1
prev <- c(0.25, 0.8)
Se.1 <- Se.2 <- 0.8
Sp.1 <- Sp.2 <- 0.95
n.sampled <- c(100, 100)
test.results <- data.frame(pp=numeric(length(prev)),
                             pn=numeric(length(prev)),
                             np=numeric(length(prev)),
                             nn=numeric(length(prev)))

for(pop in 1:length(prev)){
  n.pos <- rbinom(n.sim,n.sampled[pop],prev[pop])
  test.results$pp[pop] <- rbinom(n.sim, n.pos, Se.1*Se.2)
  rbinom(n.sim, n.sampled[pop]-n.pos, (1-Sp.1)*(1-Sp.2))
}
```

## JAGS code

```
cat(
  "model{
    # Likelihood part:
    for (i in 1:n.pop) {
      p.test.result[1,i] <-prev[i]*Se[1]*Se[2] + (1-prev[i])*Se[1]*(1-Se[2])
      p.test.result[3,i] <-prev[i]*(1-Se[1])*Se[2] + (1-prev[i])*(1-Se[1])*(1-Se[2])
      p.test.result[2,i] <-prev[i]*Se[1]*(1-Se[2]) + (1-prev[i])*Se[1]*Se[2]
      p.test.result[4,i] <-prev[i]*(1-Se[1])*(1-Se[2]) + (1-prev[i])*(1-Se[1])*Se[2]
      test.results[,i] ~dmulti(p.test.result[,i], n.sampled[i])
    }

    # Prior part:
    for (pop in 1:n.pop) {
      prev[pop] ~ dbeta(1,1)
    }
    for(test in 1:2) {
      Se[5+test] ~ dunif(0,1)
    }
  }
```

## Analysis

```
runjags.options(silent.jags=TRUE)
n.burnin <- n.sample <- 5000
results.jags <- run.jags('h-w.jags', n.chains=2, burnin=n.burnin)

## Loading required namespace: rjags

## Warning: No initial values were provided - JAGS will use
## values for all chains

## Finished running the simulation

pander(summary(results.jags))
```

Table 1: Table continues below

	Lower95	Median	Upper95	Mean	SD	Mode
<b>prev[1]</b>	0.1109	0.2135	0.337	0.2162	0.05855	0.2111
<b>prev[2]</b>	0.6232	0.7577	0.8772	0.7553	0.06486	0.7679
<b>Se[1]</b>	0.6239	0.7402	0.8462	0.7398	0.05615	0.7397
<b>Se[2]</b>	0.7347	0.8466	0.9484	0.8446	0.05529	0.8526

# Graphs

```
plot(results.jags)
```

```
## Generating plots...
```

# Exercise

- ▶ 3 tests; 1 population
- ▶ 7 parameters:
  - ▶ 3 test sensitivities
  - ▶ 3 test specificities
  - ▶ 1 prevalence
- ▶  $2^3$  combinations: 7 df in the data
  - ▶ is this identifiable?
  - ▶ are the estimates unbiased?
  - ▶ what if prevalence is very low  $\sim 1\%$ ?
  - ▶ even with 1000 individuals only  $\sim 10$  are positive
  - ▶ can't estimate Se well
  - ▶ does a biased estimate of Se bias our estimates of Sp and/or prevalence?



## Example R code

```
simulation.3.test <- function(prev,n.sampled,Se,Sp) {  
  if((length(prev)!=1)&(length(n.sampled)!=1)&(length(Se)!=1)&(length(Sp)!=1)) {  
    print("Error in parameters sent to simulation.3.test")  
    stop()  
  }  
  n.pos <- rbinom(1,n.sampled,prev)  
  n.neg <- n.sampled - n.pos  
  test.1 <- c(rbinom(n.pos,1,Se[1]), rbinom(n.neg,1,(1-Sp[1])))  
  test.2 <- c(rbinom(n.pos,1,Se[2]), rbinom(n.neg,1,(1-Sp[2])))  
  test.3 <- c(rbinom(n.pos,1,Se[3]), rbinom(n.neg,1,(1-Sp[3])))  
  test.results <- c(  
    sum(test.1 & test.2 & test.3),  
    sum(test.1 & test.2 & !test.3),  
    sum(test.1 & !test.2 & test.3),  
    sum(test.1 & !test.2 & !test.3),  
    sum(!test.1 & test.2 & test.3),  
    sum(!test.1 & test.2 & !test.3),  
    sum(!test.1 & !test.2 & test.3),  
    sum(!test.1 & !test.2 & !test.3))  
}
```

## Example R/JAGS code

```
cat(
  "model{
    # Likelihood part:
    p.test.result[1] <-prev*Se[1]*Se[2]*Se[3] + (1-prev)*(1-Se[1])*Se[2]*Se[3]
    p.test.result[2] <-prev*Se[1]*Se[2]*(1-Se[3]) + (1-prev)*Se[1]*Se[2]*(1-Se[3])
    p.test.result[3] <-prev*Se[1]*(1-Se[2])*Se[3] + (1-prev)*Se[1]*(1-Se[2])*Se[3]
    p.test.result[4] <-prev*Se[1]*(1-Se[2])*(1-Se[3]) + (1-prev)*Se[1]*(1-Se[2])*(1-Se[3])
    p.test.result[5] <-prev*(1-Se[1])*Se[2]*Se[3] + (1-prev)*(1-Se[1])*Se[2]*Se[3]
    p.test.result[6] <-prev*(1-Se[1])*Se[2]*(1-Se[3]) + (1-prev)*(1-Se[1])*Se[2]*(1-Se[3])
    p.test.result[7] <-prev*(1-Se[1])*(1-Se[2])*Se[3] + (1-prev)*(1-Se[1])*(1-Se[2])*Se[3]
    p.test.result[8] <-prev*(1-Se[1])*(1-Se[2])*(1-Se[3]) + (1-prev)*(1-Se[1])*(1-Se[2])*(1-Se[3])
    test.results ~dmulti(p.test.result, n.tested)

    # Prior part:
    prev ~ dbeta(1,1)
    for(test in 1:3) {
      Se[test] ~dbeta(1,1)
      Sp[test] ~dbeta(1,1)
    }
  }
```

# How good is our posterior?

- ▶ Eyeball posterior distribution
- ▶ Are the means (medians, modes?) close to the values used for the simulation
- ▶ If we wish to be more formal about this we would repeat the simulation-analysis cycle many (400+) times
  - ▶ this takes a long time, typically
  - ▶ which is a better (less biased) predictor: mean, median or mode
  - ▶ are the 95% Credible Intervals true 95% Confidence Intervals
  - ▶ varying the parameter values within sensible ranges

## Multiple simulations

```
set.seed(1)
n.sim <- 1000
n.pop <- 2
comparison.df <- data.frame(prev = c(runif(n.sim, 0.05, 0.95),
                                     Se.1 = runif(n.sim, 0.65, 0.95),
                                     Sp.1 = rbeta(n.sim, 21,1),
                                     Se.2 = rbeta(n.sim, 17,5),
                                     Se.2 = rbeta(n.sim, 19,3),
                                     prev.1.median = numeric(n.sim),
                                     prev.2.median = numeric(n.sim),
                                     prev.1.mean = numeric(n.sim),
                                     prev.2.mean = numeric(n.sim),
                                     Se.1.lcl = numeric(n.sim),
                                     Se.1.ucl = numeric(n.sim),
                                     Se.2.lcl = numeric(n.sim),
                                     Se.2.ucl = numeric(n.sim))

for (sim in 1:n.sim) {
  ## Simulate data
  ## ...
```

# Summary

- ▶ What is our interest?
  - ▶ which summary is better?
  - ▶ how accurate is the summary?
  - ▶ can we eliminate certain values?
  - ▶ 95% CI?
- ▶ What is the context?
  - ▶ range of possible values
  - ▶ effect of priors
  - ▶ convergence
  - ▶ sample size (see Matt)
- ▶ What if we believe that a parameter comes from a distribution, not a point value?